

Early Stopping is Nonparametric Variational Inference



Dougal Maclaurin, David Duvenaud, Ryan Adams



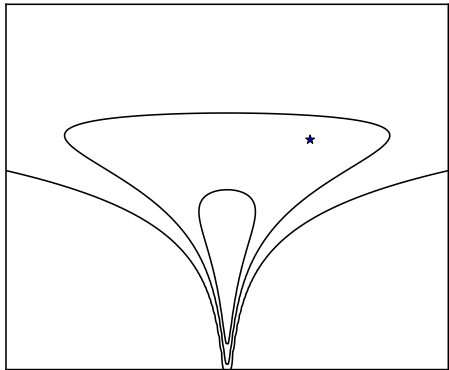
HARVARD
School of Engineering
and Applied Sciences

Good ideas always have Bayesian interpretations

Regularization	=	MAP inference
Limiting model capacity	=	Bayesian Occam's razor
Cross-validation	=	Estimating marginal likelihood
Dropout	=	Integrating out spike-and-slab
Ensembling	=	Bayes model averaging?
Early stopping	=	??

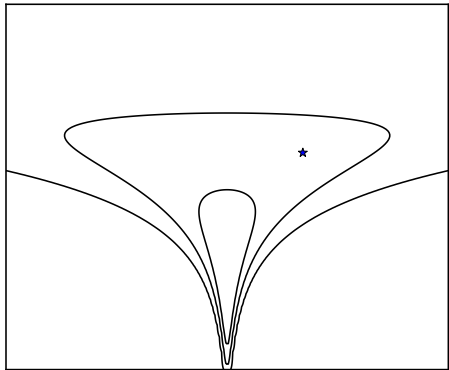
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



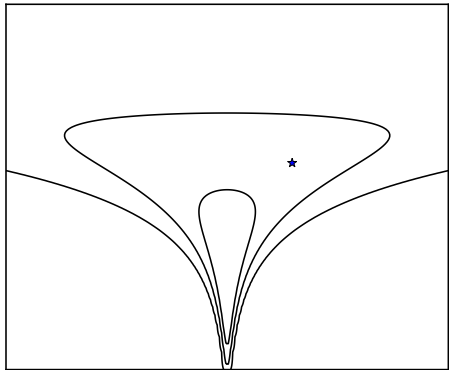
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



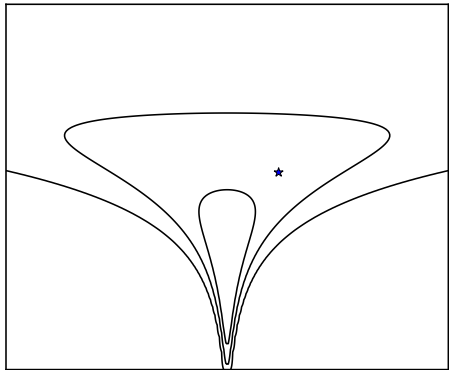
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



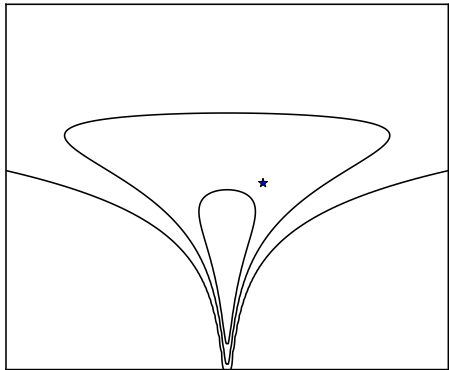
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



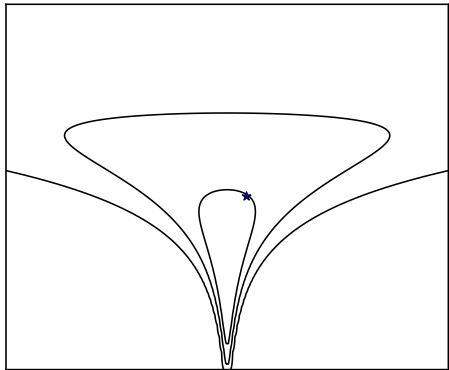
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



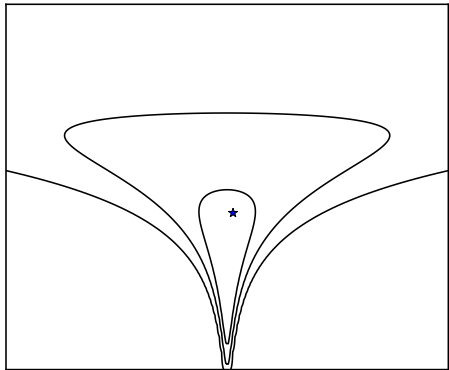
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



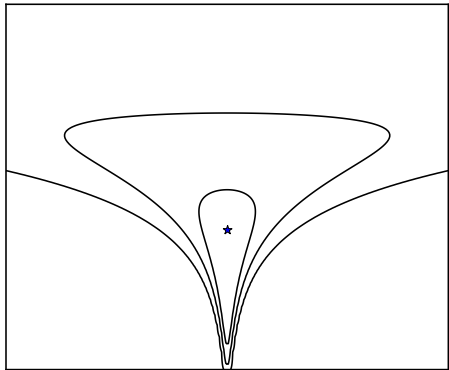
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



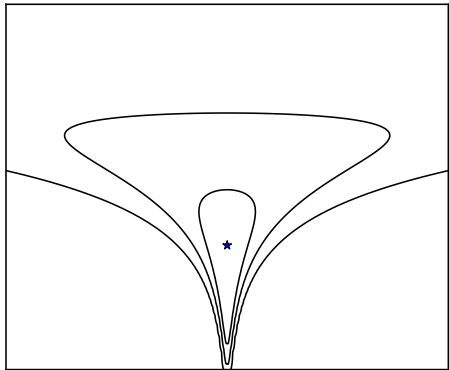
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



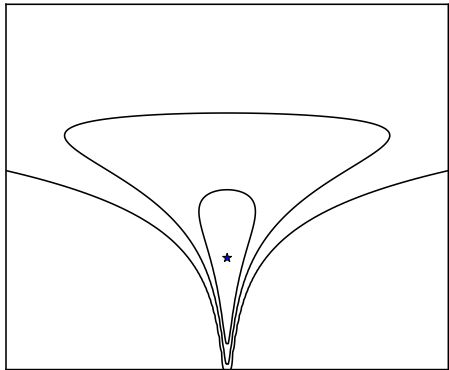
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



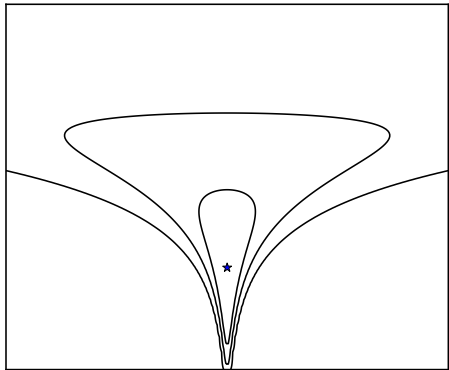
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



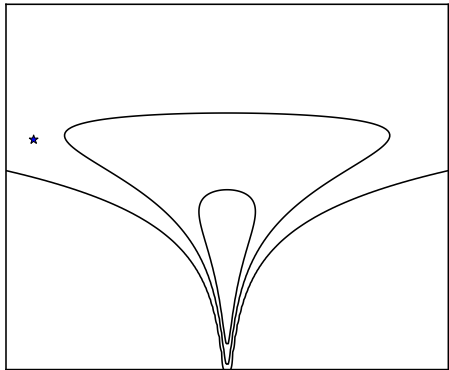
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



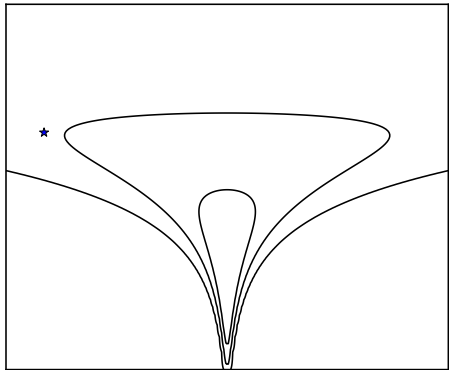
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



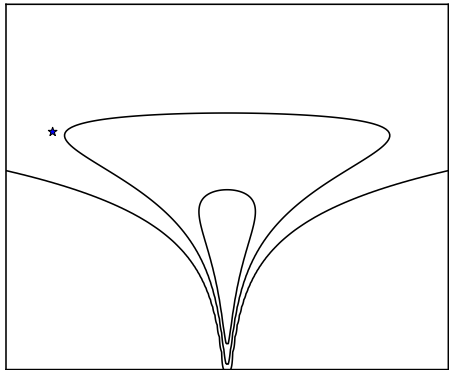
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



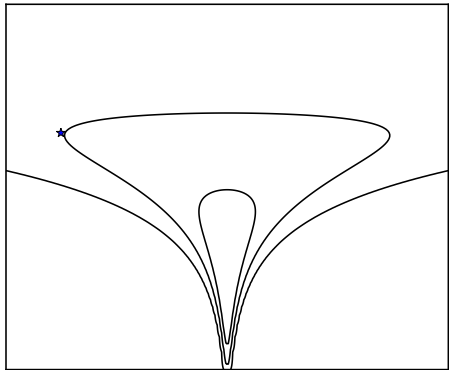
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



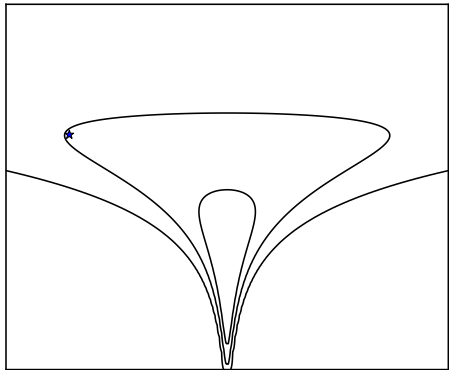
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



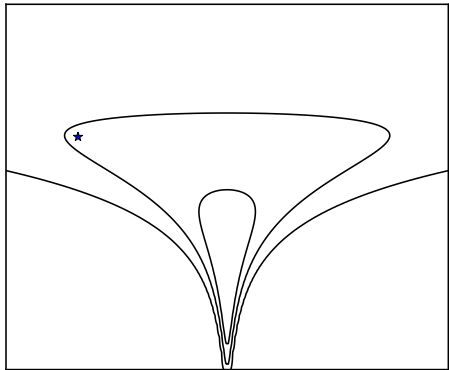
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



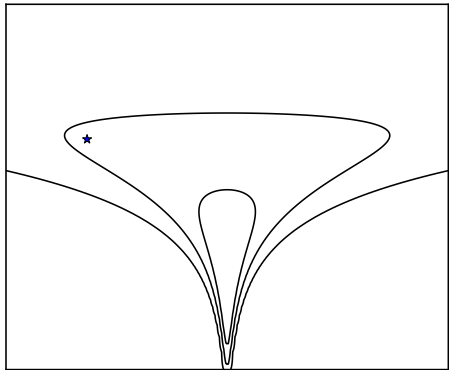
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



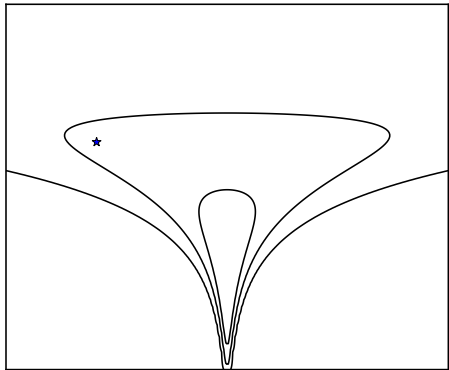
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



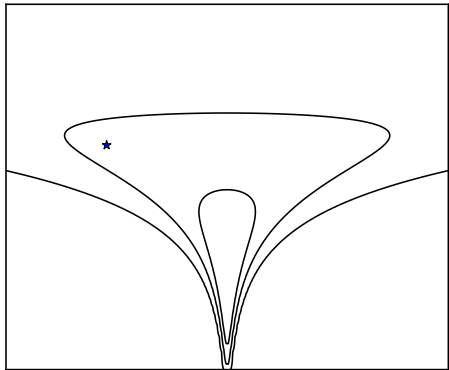
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



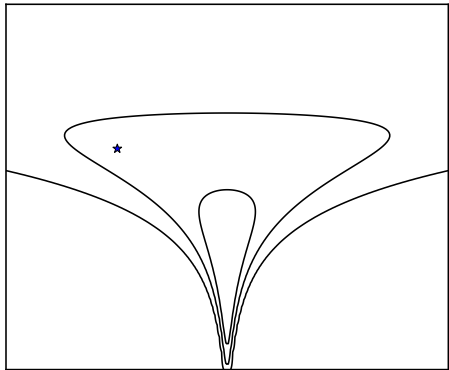
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



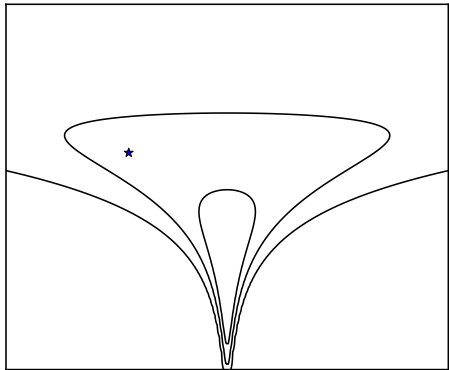
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



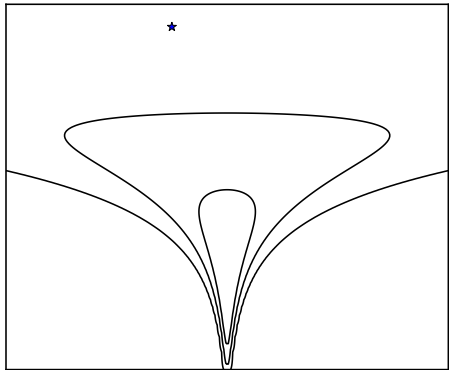
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



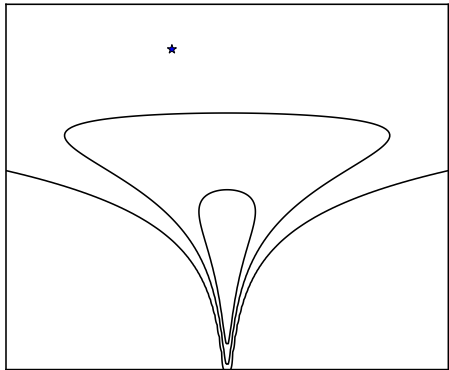
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



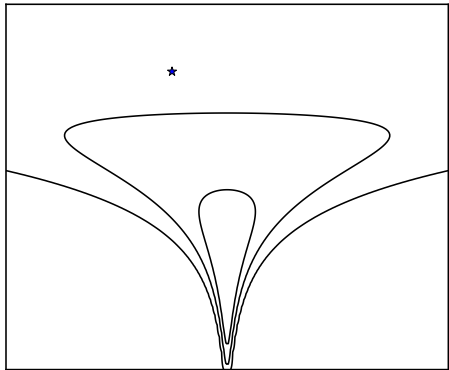
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



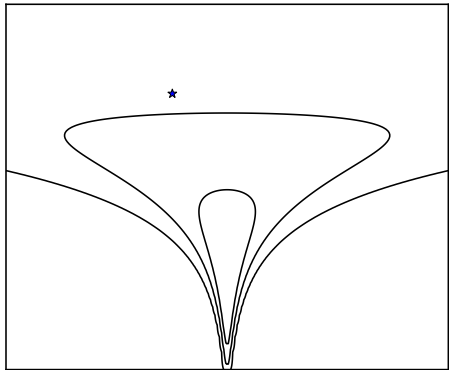
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



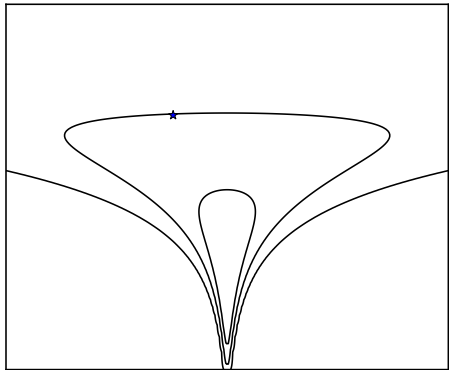
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



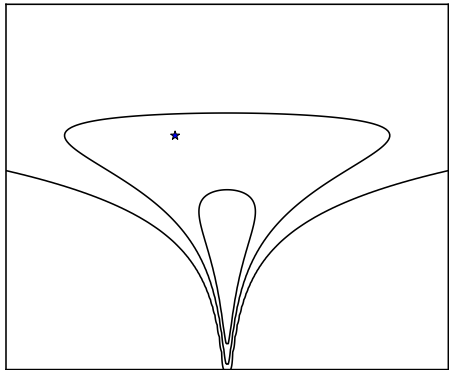
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



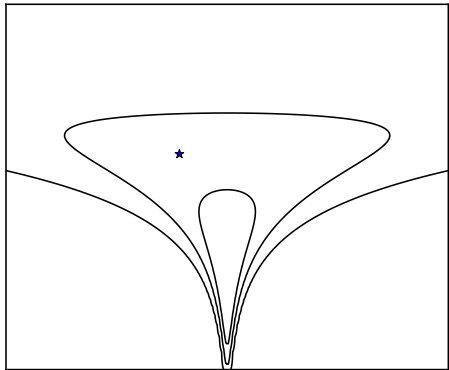
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



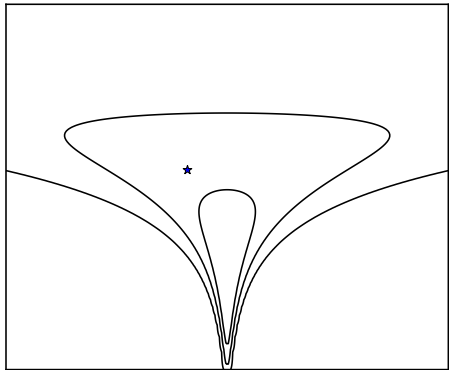
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



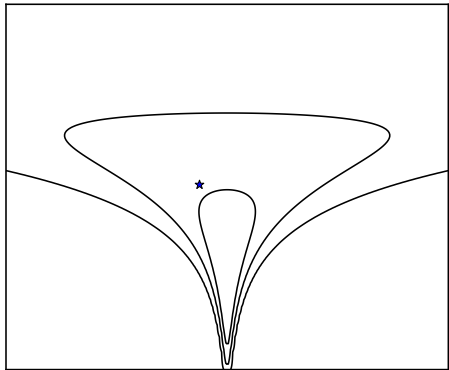
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



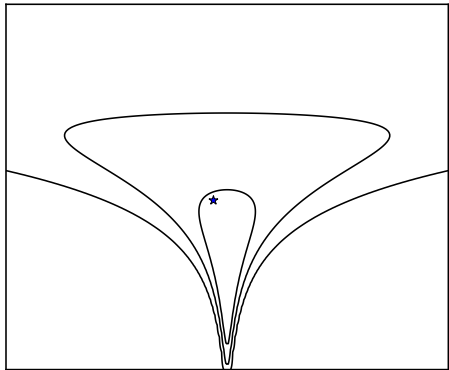
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



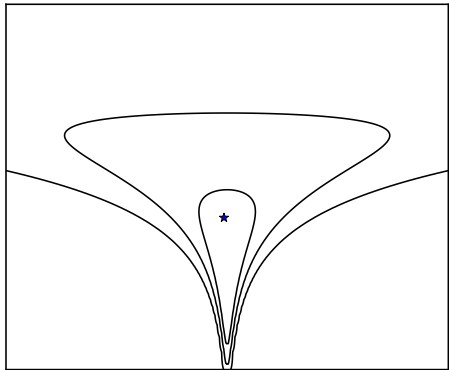
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



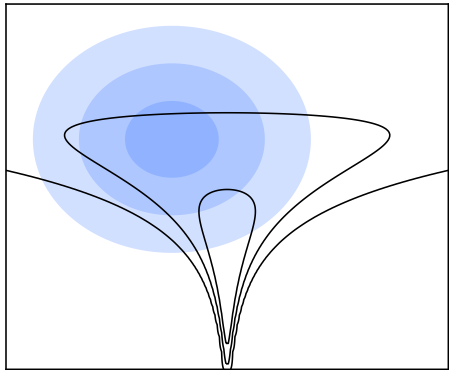
Gradient descent as a sampler

- Optimization paths start from random init, and move towards modes...



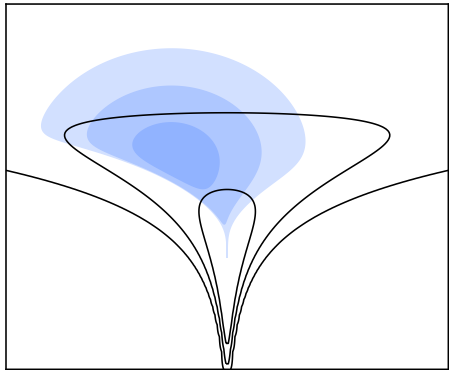
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



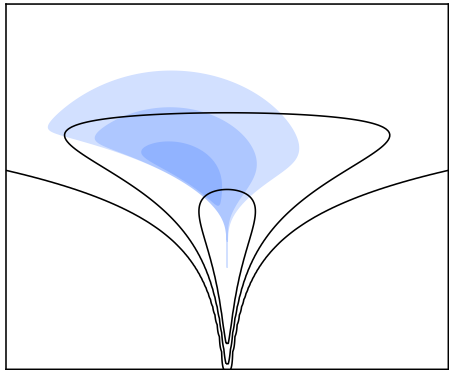
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



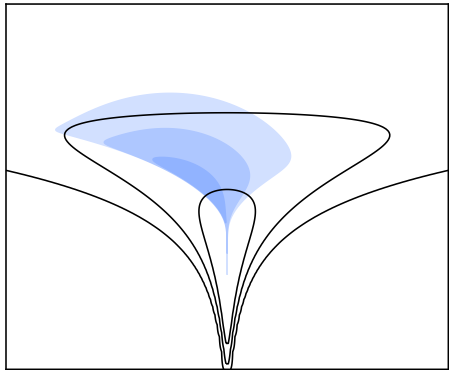
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



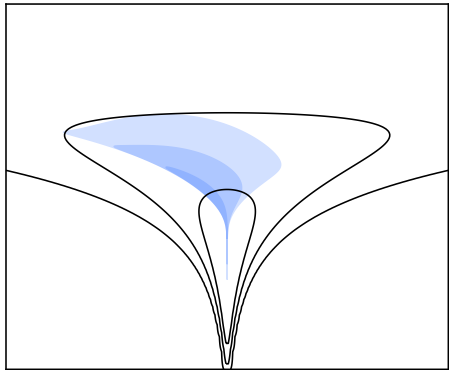
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



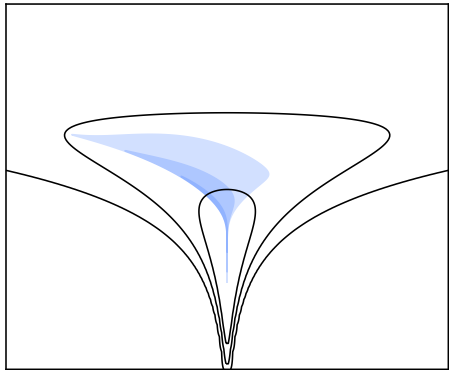
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



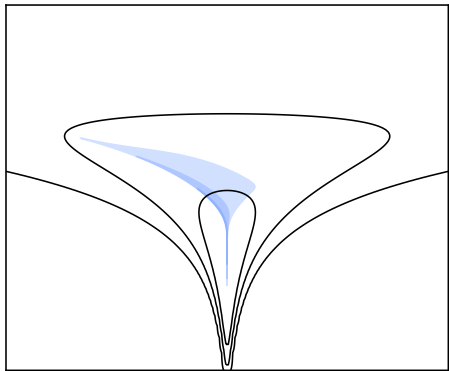
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



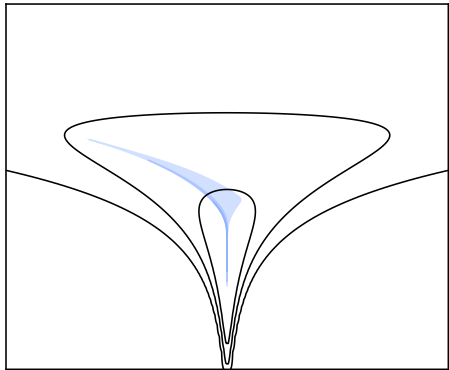
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



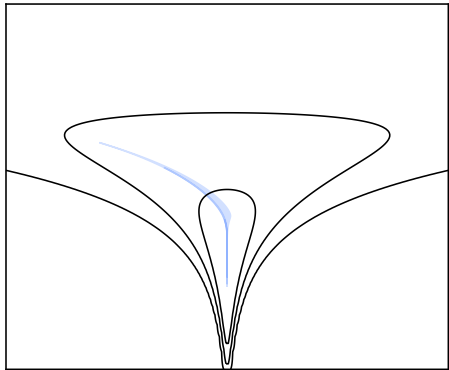
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



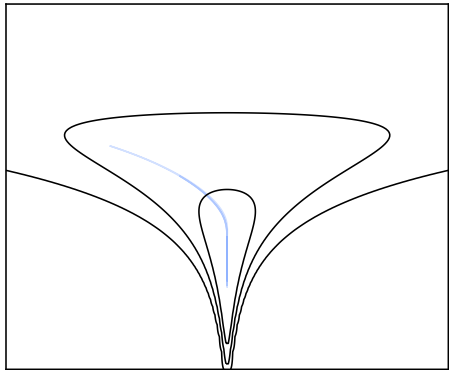
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



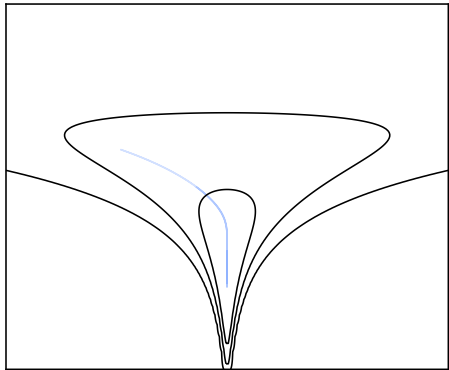
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



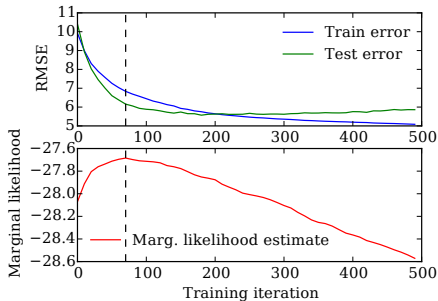
Implicit Distributions

- What about the implicit distribution of parameters after optimizing for t steps?
- Starts as a bad approximation (prior dist)
- Ends as a bad approximation (point mass)
- Ensembling = taking multiple samples from dist
- Early stopping = choosing best intermediate dist



Cross validation vs. marginal likelihood

- What if we could evaluate marginal likelihood of implicit distribution?
- Could choose all hypers to maximize marginal likelihood
- No need for cross-validation?



Variational Lower Bound

$$\log p(\mathbf{x}) \geq - \underbrace{\mathbb{E}_{q(\theta)} [-\log p(\theta, \mathbf{x})]}_{\text{Energy } E[q]} - \underbrace{\mathbb{E}_{q(\theta)} [\log q(\theta)]}_{\text{Entropy } S[q]}$$

Energy estimated from optimized objective function
(training loss is NLL):

$$\mathbb{E}_{q(\theta)} [-\log p(\theta, \mathbf{x})] \approx -\log p(\hat{\theta}_T, \mathbf{x})$$

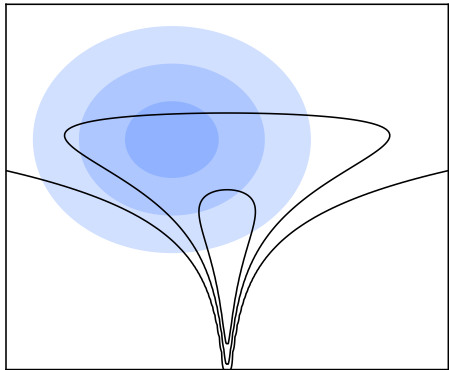
Entropy estimated by tracking change at each iteration:

$$-\mathbb{E}_{q(\theta)} [\log q(\theta)] \approx S[q_0] + \sum_{t=0}^{T-1} \log \left| J(\hat{\theta}_t) \right|$$

Using a single sample!

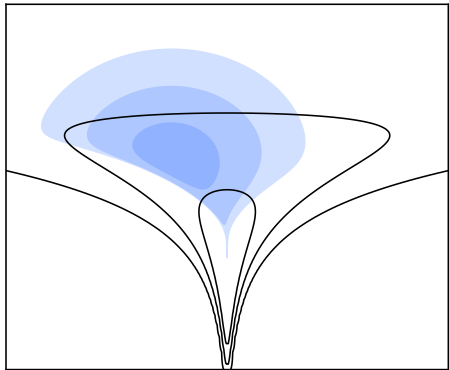
Estimating change in entropy

- Intuitively: High curvature makes entropy decrease quickly
- Can measure local curvature with Hessian
- Approximation good for small step-sizes



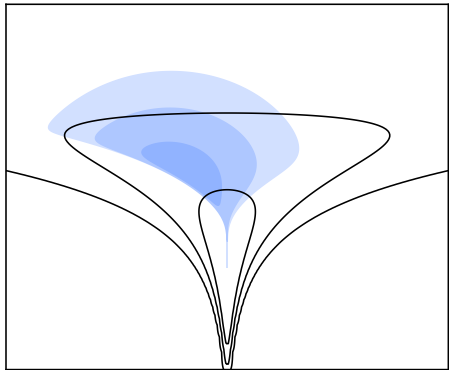
Estimating change in entropy

- Intuitively: High curvature makes entropy decrease quickly
- Can measure local curvature with Hessian
- Approximation good for small step-sizes



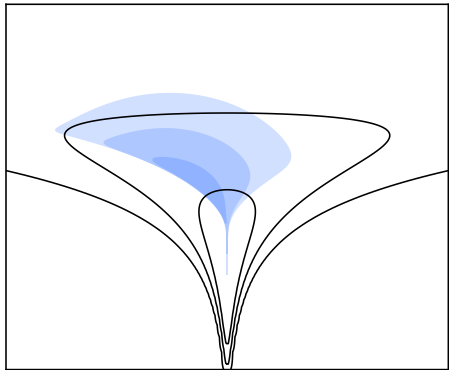
Estimating change in entropy

- Intuitively: High curvature makes entropy decrease quickly
- Can measure local curvature with Hessian
- Approximation good for small step-sizes



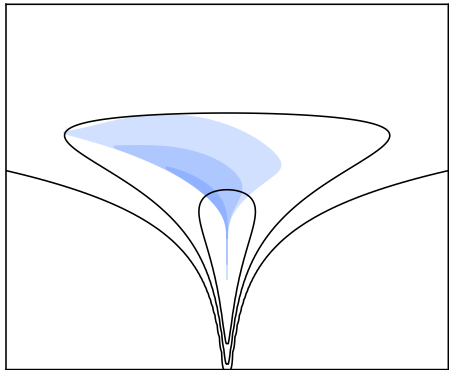
Estimating change in entropy

- Intuitively: High curvature makes entropy decrease quickly
- Can measure local curvature with Hessian
- Approximation good for small step-sizes



Estimating change in entropy

- Intuitively: High curvature makes entropy decrease quickly
- Can measure local curvature with Hessian
- Approximation good for small step-sizes



Estimating change in entropy

Volume change given by Jacobian of optimizer's operator:

$$S[q_{t+1}] - S[q_t] = \mathbb{E}_{q_t(\theta_t)} \left[\log |J(\theta_t)| \right]$$

Gradient descent update rule:

$$\theta_{t+1} = \theta_t - \alpha \nabla L(\theta),$$

Has Jacobian:

$$J(\theta_t) = I - \alpha \nabla \nabla L(\theta_t)$$

Entropy change estimated at a single sample:

$$S[q_{t+1}] - S[q_t] \approx \log |I - \alpha \nabla \nabla L(\theta_t)|$$

Final algorithm

Stochastic gradient descent

```
1: input: Weight init scale  $\sigma_0$ , step size  $\alpha$ ,  
   negative log-likelihood  $L(\theta, t)$   
2: initialize  $\theta_0 \sim \mathcal{N}(0, \sigma_0 \mathbf{I}_D)$   
3:  
4: for  $t = 1$  to  $T$  do  
5:  
6:    $\theta_t = \theta_{t-1} - \alpha \nabla L(\theta_t, t)$   
7: output sample  $\theta_T$ ,
```

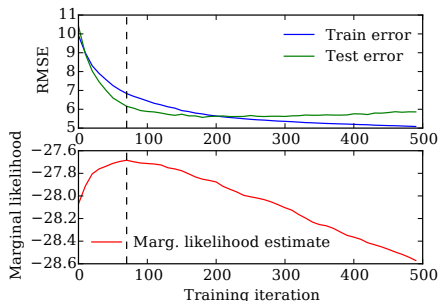
SGD with entropy estimate

```
1: input: Weight init scale  $\sigma_0$ , step size  $\alpha$ ,  
   negative log-likelihood  $L(\theta, t)$   
2: initialize  $\theta_0 \sim \mathcal{N}(0, \sigma_0 \mathbf{I}_D)$   
3: initialize  $S_0 = \frac{D}{2}(1 + \log 2\pi) + D \log \sigma_0$   
4: for  $t = 1$  to  $T$  do  
5:    $S_t = S_{t-1} + \log |\mathbf{I} - \alpha \nabla \nabla L(\theta_t, t)|$   
6:    $\theta_t = \theta_{t-1} - \alpha \nabla L(\theta_t, t)$   
7: output sample  $\theta_T$ , entropy estimate  $S_T$ 
```

- Approximate bound: $\log p(\mathbf{x}) \gtrsim -L(\theta_T) + S_T$
- Determinant is $\mathcal{O}(D^3)$
- $\mathcal{O}(D)$ Taylor approximation using Hessian-vector products
- Scales linearly in parameters and dataset size

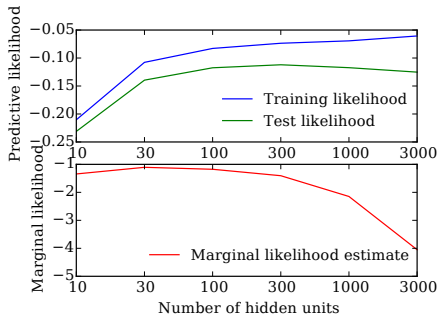
Choosing when to stop

- Neural network on the Boston housing dataset.
- SGD marginal likelihood estimate gives stopping criterion without a validation set



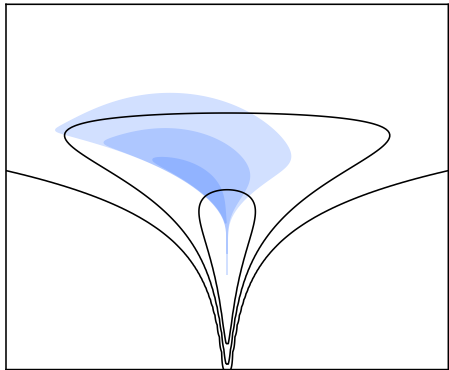
Choosing number of hidden units

- Neural net on 50000 MNIST examples
- Largest model has 2 million parameters
- Gives reasonable estimates, but cross-validation still better



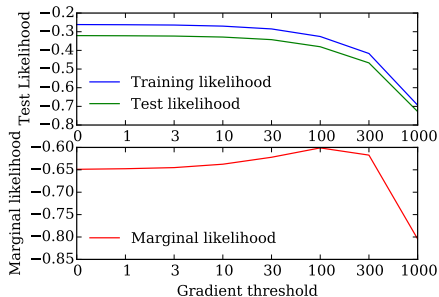
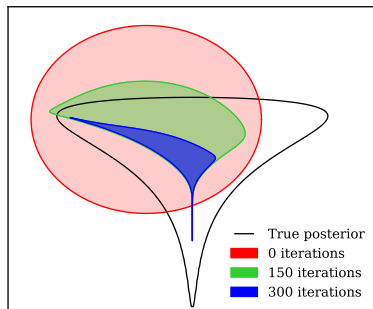
Limitations

- SGD not even trying to maximize lower bound – good approximation is by accident!
- Entropy term gets arbitrarily bad due to concentration, but true performance only gets as bad as maximum likelihood estimate



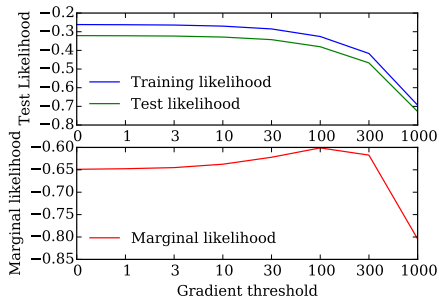
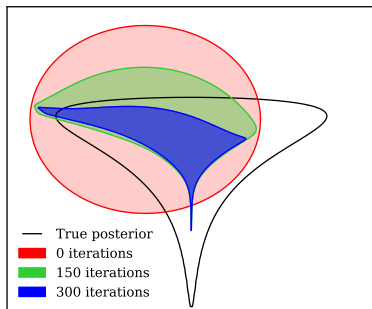
Entropy-friendly optimization

- Modified SGD to move slower near convergence, optimized new hyperparameter
- Hurts performance, but gives tighter bound
- ideally would match test likelihood



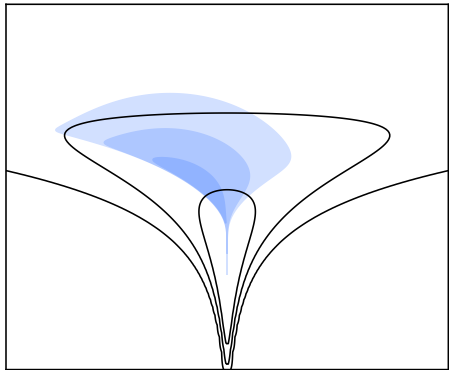
Entropy-friendly optimization

- Modified SGD to move slower near convergence, optimized new hyperparameter
- Hurts performance, but gives tighter bound
- ideally would match test likelihood



Limitations

- Irrelevant parameters can cause low entropy estimate
- No momentum - would need to estimate distribution (see Kingma & Welling, 2015)



Main Takeaways

- Optimization with random restarts implies nonparametric intermediate distributions
- Early stopping chooses among these distributions
- Ensembling samples from them
- Can scalably estimate variational lower bound on model evidence during optimization
- Another connection between practice and theory

Main Takeaways

- Optimization with random restarts implies nonparametric intermediate distributions
- Early stopping chooses among these distributions
- Ensembling samples from them
- Can scalably estimate variational lower bound on model evidence during optimization
- Another connection between practice and theory

Thanks!